

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.metrics import classification_report, confusion_matrix,
accuracy_score
```

```
aq_=pd.read_csv(r"C:\Users\siddh\Downloads\
updated_pollution_dataset.csv")
aq_
```

| | Temperature | Humidity | PM2.5 | PM10 | N02 | S02 | C0 | \ |
|------|-------------|----------|-------|-------|------|------|------|---|
| 0 | 29.8 | 59.1 | 5.2 | 17.9 | 18.9 | 9.2 | 1.72 | |
| 1 | 28.3 | 75.6 | 2.3 | 12.2 | 30.8 | 9.7 | 1.64 | |
| 2 | 23.1 | 74.7 | 26.7 | 33.8 | 24.4 | 12.6 | 1.63 | |
| 3 | 27.1 | 39.1 | 6.1 | 6.3 | 13.5 | 5.3 | 1.15 | |
| 4 | 26.5 | 70.7 | 6.9 | 16.0 | 21.9 | 5.6 | 1.01 | |
| ... | ... | ... | ... | ... | ... | ... | ... | |
| 4995 | 40.6 | 74.1 | 116.0 | 126.7 | 45.5 | 25.7 | 2.11 | |
| 4996 | 28.1 | 96.9 | 6.9 | 25.0 | 25.3 | 10.8 | 1.54 | |
| 4997 | 25.9 | 78.2 | 14.2 | 22.1 | 34.8 | 7.8 | 1.63 | |
| 4998 | 25.3 | 44.4 | 21.4 | 29.0 | 23.7 | 5.7 | 0.89 | |
| 4999 | 24.1 | 77.9 | 81.7 | 94.3 | 23.2 | 10.5 | 1.38 | |

| | Proximity_to_Industrial_Areas | Population_Density | Air Quality |
|------|-------------------------------|--------------------|-------------|
| 0 | 6.3 | 319 | Moderate |
| 1 | 6.0 | 611 | Moderate |
| 2 | 5.2 | 619 | Moderate |
| 3 | 11.1 | 551 | Good |
| 4 | 12.7 | 303 | Good |
| ... | ... | ... | ... |
| 4995 | 2.8 | 765 | Hazardous |
| 4996 | 5.7 | 709 | Moderate |
| 4997 | 9.6 | 379 | Moderate |
| 4998 | 11.6 | 241 | Good |
| 4999 | 8.3 | 461 | Moderate |

[5000 rows x 10 columns]

```
aq_.shape
```

```
(5000, 10)
```

```
aq_.describe()
```

| | Temperature | Humidity | PM2.5 | PM10 | N02 |
|-------|-------------|-------------|-------------|-------------|-------------|
| \ | | | | | |
| count | 5000.000000 | 5000.000000 | 5000.000000 | 5000.000000 | 5000.000000 |

| | | | | | |
|------|-----------|------------|------------|------------|-----------|
| mean | 30.029020 | 70.056120 | 20.142140 | 30.218360 | 26.412100 |
| std | 6.720661 | 15.863577 | 24.554546 | 27.349199 | 8.895356 |
| min | 13.400000 | 36.000000 | 0.000000 | -0.200000 | 7.400000 |
| 25% | 25.100000 | 58.300000 | 4.600000 | 12.300000 | 20.100000 |
| 50% | 29.000000 | 69.800000 | 12.000000 | 21.700000 | 25.300000 |
| 75% | 34.000000 | 80.300000 | 26.100000 | 38.100000 | 31.900000 |
| max | 58.600000 | 128.100000 | 295.000000 | 315.800000 | 64.900000 |

| | S02 | C0 | Proximity_to_Industrial_Areas \ |
|-------|-------------|-------------|---------------------------------|
| count | 5000.000000 | 5000.000000 | 5000.000000 |
| mean | 10.014820 | 1.500354 | 8.425400 |
| std | 6.750303 | 0.546027 | 3.610944 |
| min | -6.200000 | 0.650000 | 2.500000 |
| 25% | 5.100000 | 1.030000 | 5.400000 |
| 50% | 8.000000 | 1.410000 | 7.900000 |
| 75% | 13.725000 | 1.840000 | 11.100000 |
| max | 44.900000 | 3.720000 | 25.800000 |

| | Population_Density |
|-------|--------------------|
| count | 5000.000000 |
| mean | 497.423800 |
| std | 152.754084 |
| min | 188.000000 |
| 25% | 381.000000 |
| 50% | 494.000000 |
| 75% | 600.000000 |
| max | 957.000000 |

aq_.dtypes

| | |
|-------------------------------|---------|
| Temperature | float64 |
| Humidity | float64 |
| PM2.5 | float64 |
| PM10 | float64 |
| N02 | float64 |
| S02 | float64 |
| C0 | float64 |
| Proximity_to_Industrial_Areas | float64 |
| Population_Density | int64 |
| Air Quality | object |
| dtype: | object |

aq_.head(5)

| | Temperature | Humidity | PM2.5 | PM10 | N02 | S02 | C0 | \ |
|---|-------------|----------|-------|------|------|------|------|---|
| 0 | 29.8 | 59.1 | 5.2 | 17.9 | 18.9 | 9.2 | 1.72 | |
| 1 | 28.3 | 75.6 | 2.3 | 12.2 | 30.8 | 9.7 | 1.64 | |
| 2 | 23.1 | 74.7 | 26.7 | 33.8 | 24.4 | 12.6 | 1.63 | |
| 3 | 27.1 | 39.1 | 6.1 | 6.3 | 13.5 | 5.3 | 1.15 | |
| 4 | 26.5 | 70.7 | 6.9 | 16.0 | 21.9 | 5.6 | 1.01 | |

| | Proximity_to_Industrial_Areas | Population_Density | Air Quality |
|---|-------------------------------|--------------------|-------------|
| 0 | 6.3 | 319 | Moderate |
| 1 | 6.0 | 611 | Moderate |
| 2 | 5.2 | 619 | Moderate |
| 3 | 11.1 | 551 | Good |
| 4 | 12.7 | 303 | Good |

```
aq_.tail(5)
```

| | Temperature | Humidity | PM2.5 | PM10 | N02 | S02 | CO | \ |
|------|-------------|----------|-------|-------|------|------|------|---|
| 4995 | 40.6 | 74.1 | 116.0 | 126.7 | 45.5 | 25.7 | 2.11 | |
| 4996 | 28.1 | 96.9 | 6.9 | 25.0 | 25.3 | 10.8 | 1.54 | |
| 4997 | 25.9 | 78.2 | 14.2 | 22.1 | 34.8 | 7.8 | 1.63 | |
| 4998 | 25.3 | 44.4 | 21.4 | 29.0 | 23.7 | 5.7 | 0.89 | |
| 4999 | 24.1 | 77.9 | 81.7 | 94.3 | 23.2 | 10.5 | 1.38 | |

| | Proximity_to_Industrial_Areas | Population_Density | Air Quality |
|------|-------------------------------|--------------------|-------------|
| 4995 | 2.8 | 765 | 1 |
| 4996 | 5.7 | 709 | 2 |
| 4997 | 9.6 | 379 | 2 |
| 4998 | 11.6 | 241 | 0 |
| 4999 | 8.3 | 461 | 2 |

```
aq_.columns
```

```
Index(['Temperature', 'Humidity', 'PM2.5', 'PM10', 'N02', 'S02', 'CO',
      'Proximity_to_Industrial_Areas', 'Population_Density', 'Air
      Quality'],
      dtype='object')
```

```
aq_.isna()
```

[illegible]

| | Proximity_to_Industrial_Areas | Population_Density | Air Quality |
|------|-------------------------------|--------------------|-------------|
| 0 | False | False | False |
| 1 | False | False | False |
| 2 | False | False | False |
| 3 | False | False | False |
| 4 | False | False | False |
| ... | ... | ... | ... |
| 4995 | False | False | False |
| 4996 | False | False | False |
| 4997 | False | False | False |
| 4998 | False | False | False |
| 4999 | False | False | False |

[5000 rows x 10 columns]

aq_.isnull().sum()

| | |
|-------------------------------|---|
| Temperature | 0 |
| Humidity | 0 |
| PM2.5 | 0 |
| PM10 | 0 |
| N02 | 0 |
| S02 | 0 |
| C0 | 0 |
| Proximity_to_Industrial_Areas | 0 |
| Population_Density | 0 |
| Air Quality | 0 |

dtype: int64

aq_.info()

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 5000 entries, 0 to 4999

Data columns (total 10 columns):

| # | Column | Non-Null Count | Dtype |
|-----|-------------------------------|----------------|---------|
| --- | ----- | ----- | ----- |
| 0 | Temperature | 5000 non-null | float64 |
| 1 | Humidity | 5000 non-null | float64 |
| 2 | PM2.5 | 5000 non-null | float64 |
| 3 | PM10 | 5000 non-null | float64 |
| 4 | N02 | 5000 non-null | float64 |
| 5 | S02 | 5000 non-null | float64 |
| 6 | C0 | 5000 non-null | float64 |
| 7 | Proximity_to_Industrial_Areas | 5000 non-null | float64 |
| 8 | Population_Density | 5000 non-null | int64 |
| 9 | Air Quality | 5000 non-null | object |

dtypes: float64(8), int64(1), object(1)

memory usage: 390.8+ KB

TO FIND THE CATAEGORICAL VALUES

```
categorical_features=[ i for i in aq_.columns if aq_[i].dtype == "Air  
Quality"]  
aq_[categorical_features]
```

Empty DataFrame

Columns: []

Index: [0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50, 51, 52, 53, 54, 55, 56, 57, 58, 59, 60, 61, 62, 63, 64, 65, 66, 67, 68, 69, 70, 71, 72, 73, 74, 75, 76, 77, 78, 79, 80, 81, 82, 83, 84, 85, 86, 87, 88, 89, 90, 91, 92, 93, 94, 95, 96, 97, 98, 99, ...]

[5000 rows x 0 columns]

```
label_encoder = LabelEncoder()
```

```
aq_['Air Quality'] = label_encoder.fit_transform(aq_['Air Quality'])  
aq_
```

| | Temperature | Humidity | PM2.5 | PM10 | N02 | S02 | C0 | \ |
|------|-------------|----------|-------|-------|------|------|------|---|
| 0 | 29.8 | 59.1 | 5.2 | 17.9 | 18.9 | 9.2 | 1.72 | |
| 1 | 28.3 | 75.6 | 2.3 | 12.2 | 30.8 | 9.7 | 1.64 | |
| 2 | 23.1 | 74.7 | 26.7 | 33.8 | 24.4 | 12.6 | 1.63 | |
| 3 | 27.1 | 39.1 | 6.1 | 6.3 | 13.5 | 5.3 | 1.15 | |
| 4 | 26.5 | 70.7 | 6.9 | 16.0 | 21.9 | 5.6 | 1.01 | |
| ... | ... | ... | ... | ... | ... | ... | ... | |
| 4995 | 40.6 | 74.1 | 116.0 | 126.7 | 45.5 | 25.7 | 2.11 | |
| 4996 | 28.1 | 96.9 | 6.9 | 25.0 | 25.3 | 10.8 | 1.54 | |
| 4997 | 25.9 | 78.2 | 14.2 | 22.1 | 34.8 | 7.8 | 1.63 | |
| 4998 | 25.3 | 44.4 | 21.4 | 29.0 | 23.7 | 5.7 | 0.89 | |
| 4999 | 24.1 | 77.9 | 81.7 | 94.3 | 23.2 | 10.5 | 1.38 | |

| | Proximity_to_Industrial_Areas | Population_Density | Air Quality |
|------|-------------------------------|--------------------|-------------|
| 0 | 6.3 | 319 | 2 |
| 1 | 6.0 | 611 | 2 |
| 2 | 5.2 | 619 | 2 |
| 3 | 11.1 | 551 | 0 |
| 4 | 12.7 | 303 | 0 |
| ... | ... | ... | ... |
| 4995 | 2.8 | 765 | 1 |
| 4996 | 5.7 | 709 | 2 |
| 4997 | 9.6 | 379 | 2 |
| 4998 | 11.6 | 241 | 0 |
| 4999 | 8.3 | 461 | 2 |


[5000 rows x 10 columns]

```
print(aq_['Air Quality'].unique())
```

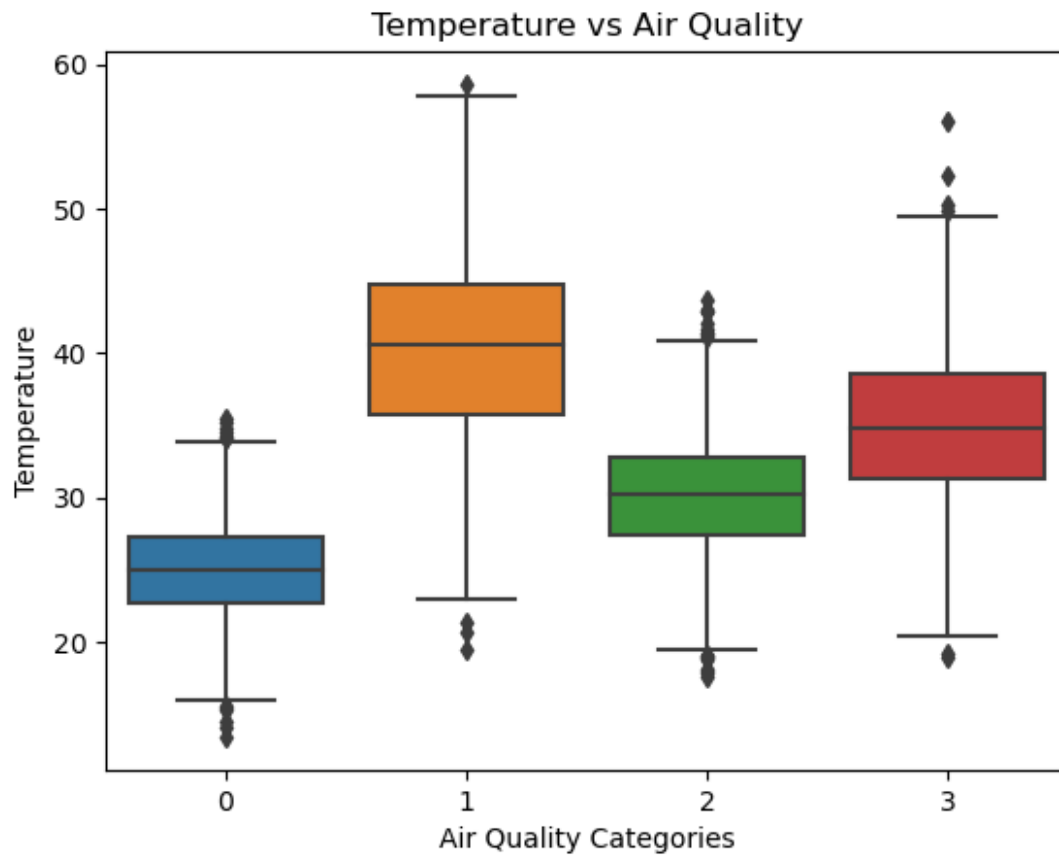
```
[2 0 1 3]
aq_.index
RangeIndex(start=0, stop=5000, step=1)
```

EDA CHART


```
sns.countplot(x='Air Quality', data=aq_)
plt.title('Distribution of Air Quality')
plt.xlabel('Air Quality Categories')
plt.ylabel('Count')
plt.show();
```



```
sns.boxplot(x='Air Quality', y='Temperature', data=aq_)
plt.title('Temperature vs Air Quality')
plt.xlabel('Air Quality Categories')
plt.ylabel('Temperature')
plt.show();
```



```
sns.scatterplot(  
    x='Proximity_to_Industrial_Areas',  
    y='Population_Density',  
    hue='Air Quality',  
    palette='coolwarm',  
    data=aq_  
)  
plt.show;
```



LOGISTIC REGRESSION MODEL

```
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

# Define features and target
X = aq_.drop(columns=['Air Quality'])
y = aq_['Air Quality']
```

to identify the x test and x test columns and values

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.3, random_state=42)

X_train
```


| | Temperature | Humidity | PM2.5 | PM10 | N02 | S02 | C0 | \ |
|-------------------------|-------------------------------|----------|-------|--------------------|------|------|------|---|
| 1840 | 39.1 | 82.9 | 56.7 | 69.3 | 39.7 | 8.2 | 2.00 | |
| 2115 | 22.1 | 83.0 | 0.7 | 9.3 | 32.5 | 11.4 | 1.60 | |
| 4437 | 34.5 | 77.7 | 6.2 | 7.1 | 25.6 | 8.9 | 1.71 | |
| 1146 | 23.8 | 65.2 | 9.9 | 14.2 | 14.6 | 4.9 | 0.97 | |
| 2486 | 34.6 | 81.4 | 21.9 | 32.1 | 29.6 | 1.9 | 1.19 | |
| ... | ... | ... | ... | ... | ... | ... | ... | |
| 4426 | 30.0 | 59.7 | 59.3 | 68.8 | 22.1 | 7.4 | 1.39 | |
| 466 | 27.7 | 54.0 | 17.6 | 23.2 | 24.6 | 11.1 | 1.27 | |
| 3092 | 24.2 | 67.6 | 19.4 | 35.1 | 24.2 | 9.6 | 1.91 | |
| 3772 | 21.4 | 76.4 | 15.3 | 22.1 | 12.4 | 3.4 | 0.96 | |
| 860 | 34.4 | 103.1 | 29.2 | 42.2 | 54.4 | 23.6 | 3.37 | |
| | | | | | | | | |
| | Proximity_to_Industrial_Areas | | | Population_Density | | | | |
| 1840 | | | | 3.6 | | | | |
| 2115 | | | | 6.3 | | | | |
| 4437 | | | | 8.1 | | | | |
| 1146 | | | | 10.7 | | | | |
| 2486 | | | | 6.9 | | | | |
| ... | | | | ... | | | | |
| 4426 | | | | 5.1 | | | | |
| 466 | | | | 7.8 | | | | |
| 3092 | | | | 6.2 | | | | |
| 3772 | | | | 16.3 | | | | |
| 860 | | | | 3.7 | | | | |
| | | | | | | | | |
| [3500 rows x 9 columns] | | | | | | | | |
| X_test | | | | | | | | |
| | Temperature | Humidity | PM2.5 | PM10 | N02 | S02 | C0 | \ |
| 1501 | 31.3 | 73.4 | 72.8 | 88.8 | 54.1 | 17.1 | 2.54 | |
| 2586 | 24.0 | 68.8 | 0.8 | 8.6 | 20.9 | 2.4 | 1.16 | |
| 2653 | 25.4 | 78.6 | 4.1 | 6.7 | 21.3 | 9.1 | 1.40 | |
| 1055 | 30.7 | 102.0 | 7.9 | 26.1 | 41.6 | 26.1 | 2.31 | |
| 705 | 26.5 | 54.1 | 12.3 | 15.3 | 20.3 | 5.3 | 0.99 | |
| ... | ... | ... | ... | ... | ... | ... | ... | |
| 3563 | 20.9 | 58.2 | 17.5 | 23.6 | 18.8 | 3.9 | 0.94 | |
| 1538 | 20.6 | 88.6 | 47.9 | 57.9 | 28.9 | 12.3 | 1.94 | |
| 1837 | 34.7 | 75.8 | 2.7 | 6.4 | 33.2 | 6.3 | 1.50 | |
| 2380 | 22.2 | 98.7 | 13.2 | 35.8 | 24.0 | 12.6 | 1.54 | |
| 1912 | 22.7 | 73.6 | 8.2 | 13.6 | 21.8 | 5.2 | 1.14 | |
| | | | | | | | | |
| | Proximity_to_Industrial_Areas | | | Population_Density | | | | |
| 1501 | | | | 2.6 | | | | |
| 2586 | | | | 10.5 | | | | |
| 2653 | | | | 10.1 | | | | |
| 1055 | | | | 4.8 | | | | |
| 705 | | | | 12.7 | | | | |
| ... | | | | ... | | | | |

| | | |
|------|------|-----|
| 3563 | 10.9 | 420 |
| 1538 | 5.4 | 412 |
| 1837 | 7.8 | 335 |
| 2380 | 3.8 | 471 |
| 1912 | 12.2 | 472 |

[1500 rows x 9 columns]

to identify the y test and y test columns and values

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.3, random_state=42)
```

y_train

| | |
|------|---|
| 1840 | 3 |
| 2115 | 2 |
| 4437 | 2 |
| 1146 | 0 |
| 2486 | 2 |

| | |
|------|----|
| | .. |
| 4426 | 2 |
| 466 | 2 |
| 3092 | 3 |
| 3772 | 0 |
| 860 | 1 |

Name: Air Quality, Length: 3500, dtype: int32

y_test

| | |
|------|---|
| 1501 | 1 |
| 2586 | 0 |
| 2653 | 2 |
| 1055 | 1 |
| 705 | 0 |

| | |
|------|----|
| | .. |
| 3563 | 0 |
| 1538 | 2 |
| 1837 | 2 |
| 2380 | 3 |
| 1912 | 0 |

Name: Air Quality, Length: 1500, dtype: int32

Split dataset (70% training, 30% testing)

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.3, random_state=42)
```

```

# Standardize features (important for regularization)
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Model with regularization (L2 penalty)
model = LogisticRegression(penalty='l2', C=0.1)

# Train the model
model.fit(X_train, y_train)

# Evaluate performance
y_train_pred = model.predict(X_train)
y_test_pred = model.predict(X_test)

train_acc = accuracy_score(y_train, y_train_pred)
test_acc = accuracy_score(y_test, y_test_pred)

print(f"Training Accuracy: {train_acc:.2f}")
print(f"Testing Accuracy: {test_acc:.2f}")

# Corrected accuracy calculation
print(" Accuracy:", round(test_acc * 100, 2))

Training Accuracy: 0.94
Testing Accuracy: 0.93
Accuracy: 93.2

from sklearn.metrics import classification_report

# Generate classification report
report = classification_report(y_test, y_test_pred)

print(report)

```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.99 | 1.00 | 1.00 | 618 |
| 1 | 0.88 | 0.78 | 0.83 | 148 |
| 2 | 0.94 | 0.94 | 0.94 | 452 |
| 3 | 0.82 | 0.85 | 0.83 | 282 |
| accuracy | | | 0.93 | 1500 |
| macro avg | 0.91 | 0.89 | 0.90 | 1500 |
| weighted avg | 0.93 | 0.93 | 0.93 | 1500 |

```

print(y_test_pred[:10]) # Shows first 10 predictions
print(y_test[:10].values) # Shows actual values for comparison

```

```
[1 0 2 1 0 1 0 0 0 2]  
[1 0 2 1 0 1 0 0 0 2]
```